

Nikhil Reddy Chinthakunta

Senior Data Engineer · Databricks & Delta Lake · PySpark · Medallion Architecture · Healthcare · HIPAA
New York, NY · +1 (216) 235-6712 · chinthakuntanikhilreddy@gmail.com · [LinkedIn](#)

PROFESSIONAL SUMMARY

Senior Data Engineer with 7+ years of experience building and operating production data platforms on the Databricks Lakehouse designing, building, and maintaining Bronze, Silver, and Gold medallion pipelines that power member profiles, clinical signals, engagement data, and decision intelligence use cases. Strong SQL and Python with hands-on Spark-based pipeline development. Deep Databricks experience including Feature Store, Unity Catalog, Delta Live Tables, and Databricks Workflows. Builds batch and streaming ingestion pipelines using Databricks, Spark Structured Streaming, Delta Lake, Kafka, and event-driven data architectures. Implements data quality validation, reconciliation, monitoring, alerting, and testing controls ensuring trusted production-ready data assets. Designs scalable data models balancing usability, governance, performance, and long-term maintainability. Healthcare domain expertise HIPAA-compliant PHI handling, clinical data, claims, EHR, and social determinants of health at IBM. Experience with Azure Data Factory, Azure Data Lake Storage, and Azure Event Hubs. Familiarity with ML, reinforcement learning, and decision intelligence platforms. Builds with Claude Code and GitHub Copilot to work faster and at higher quality. Databricks Certified Data Engineer Associate.

TECHNICAL SKILLS

Core Stack: Python (production) · SQL (strong) · dbt (production - models, tests, monitoring, schema versioning, CI/CD) · Snowflake (production) · Airflow (orchestration) · ELT pipeline design

Databricks Platform: Databricks (deep production) · Delta Lake · Delta Live Tables · Databricks Feature Store · Unity Catalog · Databricks Workflows · Lakeflow Connect · medallion architecture (Bronze/Silver/Gold)

Spark & Processing: Apache Spark (production) · PySpark · Spark Structured Streaming · real-time feature engineering · batch and streaming pipeline design · Spark workload tuning · partitioning strategies

Streaming & Event-Driven: Kafka (production) · event-driven data architectures · Azure Event Hubs · real-time ingestion · streaming feature engineering · near-real-time decisioning data · batch and real-time patterns

Healthcare Domain: HIPAA-compliant PHI handling (production - IBM) · clinical data · EHR · claims data · member profiles · engagement data · socioeconomic/SDOH data · ICD-10/CPT/NDC/NPI · CMS/CDC adjacent

Data Quality & Governance: Data quality validation · reconciliation · monitoring · alerting · testing controls · Great Expectations (familiar) · Monte Carlo (familiar) · data lineage · security · auditability · regulatory compliance

Azure & Cloud: Azure Data Factory (production) · Azure Data Lake Storage (production) · Azure Event Hubs · Azure Databricks · AWS (S3, Glue, EMR) · cloud-native distributed data processing

ML & Decision Intelligence: Databricks Feature Store · Gold-layer feature tables · model training data · model scoring · reinforcement learning data pipelines · decision intelligence use cases · reusable data products

AI Tools: Claude Code (daily production - builds faster and at higher quality) · GitHub Copilot · AI-assisted data engineering ·

PROFESSIONAL EXPERIENCE

JPMorgan Chase & Co.

Senior Data Engineer

Apr 2025 - Present
(Remote) New York, NY

- Designed and maintained Bronze, Silver, and Gold lakehouse pipelines on Databricks built ELT pipelines moving data from Salesforce and third-party financial systems into Delta Lake translating financial stakeholder data needs into technical requirements.
- Built and maintained canonical dbt models and Gold-layer feature tables as the foundation for financial decision-making applied consistent schemas, well-tested transformation logic, and version control treating financial data models as production software.
- Developed deep Databricks and columnar data warehouse expertise applied partitioning, clustering, and Spark workload tuning reducing query execution times by 30% managing warehouse cost and performance across large complex financial datasets.
- Supported GTM, Finance, Product, and Engineering stakeholders across departments standardized KPI definitions, eliminated inconsistent reporting, and delivered a unified semantic layer as the platform-level single source of truth for financial metrics.
- Built scalable financial data systems and reusable data products reduced shadow tables and one-off datasets proactively serving team data needs freeing GTM, Finance, and Product teams to move faster with governed trustworthy data assets.
- Built the data foundation for AI and decision intelligence initiatives structured financial data models with consistent schemas optimized for ML model training, model scoring, reinforcement learning, and LLM-powered analytics using Databricks Feature Store.
- Translated complex financial requirements into scoped buildable data solutions diagnosed and resolved pipeline failures and production incidents delivering clear actionable insights for technical and non-technical stakeholders maintaining reliable platform operations.
- Applied Claude Code, GitHub Copilot, and Cursor as daily AI fluency tools accelerated dbt development, SQL optimization, and automated testing building financial data infrastructure faster and at higher quality.
- Reduced pipeline processing latency by 40% and reporting defects by 35% applied dbt model optimization, rigorous testing, data quality validation, monitoring, and observability frameworks ensuring financial datasets were trusted by both human and AI consumers.

IBM

Senior Data Engineer

Jan 2023 - Mar 2025
Plano, Texas

- Designed, built, and maintained Bronze, Silver, and Gold lakehouse pipelines on Databricks powered member profiles, clinical signals, healthcare engagement data, and decision intelligence use cases with Delta Lake and medallion architecture.
- Built batch and streaming ingestion pipelines using Databricks, Spark Structured Streaming, Kafka, and event-driven architectures applied Azure Data Factory, ADLS, and Azure Event Hubs for cloud-native healthcare data ingestion.
- Implemented data quality validation, reconciliation, monitoring, alerting, and testing controls ensured trusted, production-ready healthcare data assets with HIPAA-compliant PHI handling, data lineage, security, and auditability.

- Developed Gold-layer feature tables using Databricks Feature Store built reusable data products supporting model training, model scoring, reinforcement learning, and real-time decisioning across the NBA platform ecosystem.
- Applied Unity Catalog for data governance and lineage tracking managed schemas, permissions, and metadata across multi-workspace Databricks environments with compliance-aware engineering practices.
- Tuned Spark workloads, storage strategies, partitioning schemes, and query performance improved efficiency, scalability, and cost management reducing pipeline processing latency by 40%.
- Designed scalable healthcare data models balancing usability, governance, performance, and long-term maintainability integrated clinical, claims, EHR, member behavioral, engagement, and SDOH datasets.
- Diagnosed and resolved pipeline failures, data quality issues, and production incidents maintained reliable platform operations reducing data defects by 35% through rigorous testing and observability frameworks.
- Built with Claude Code and GitHub Copilot as daily production tools worked faster and at higher quality partnering closely with Data Science, AI Engineering, and Platform Engineering teams.

Infosys

Aug 2017 - Jul 2021

Data Engineer

Hyderabad, India

- Designed and maintained Bronze, Silver, and Gold lakehouse pipelines on Databricks to powered financial member profiles, engagement data, and decision intelligence use cases with Delta Lake and medallion architecture.
- Built batch and streaming ingestion pipelines using Spark, Kafka, and event-driven architectures implemented data quality validation, reconciliation, monitoring, and alerting ensuring trusted production-ready financial data assets.
- Developed Gold-layer feature tables supporting model training, scoring, and real-time decisioning tuned Spark workloads and partitioning strategies improving efficiency and cost management reducing errors by 30%.
- Applied data governance standards for lineage, security, PHI-handling, auditability, and regulatory compliance partnered with Data Science and AI Engineering teams ensuring high-quality data across the platform ecosystem.
- Built with GitHub Copilot diagnosed and resolved pipeline failures and production incidents maintaining reliable platform operations with strong communication across engineering and business teams.

EDUCATION

Belhaven University | *Phd., Data Science*

Trine University | *M.S., Information Studies*

Sphoorthy Engineering College | *B.S., Computer Science and Engineering*